

Lyte — Decision Intelligence

Multi-model AI routing and portfolio-wide signal observability.

Manifest snapshot · 2026-04-19


POSITIONING

Lyte is the decision intelligence layer of the SZL ecosystem — it ingests signals from every platform, routes work across an LLM model fleet with cost and policy awareness, and exposes the operational state of the business as a single observable surface.

CAPABILITY SNAPSHOT

TOTAL CAPABILITIES 5	READINESS SCORE 73%	LIVE 1	WORKING DEMO 3	PARTIAL / AT RISK 1
-------------------------	------------------------	-----------	-------------------	------------------------

STATUS MIX



TOP PROOF POINTS

CAP-024 · PRISM framework dashboard — Pulse, Risk, Intelligence, Signals, Motion

All five PRISM dimensions render with seeded signals; scoring pipeline functional; approval queue interactive
status: Working Demo · module: artifacts/szl-holdings/src/ (Lyte workspace), artifacts/lyte-command-center/src/

CAP-025 · Live intelligence feeds — BLS unemployment, GitHub Trending, TechCrunch RSS, The Verge RSS

All four feeds confirmed live — no API key required; tested in mock-stub-placeholder-register.md
status: Live · module: artifacts/api-server/src/lib/intelligence-feeds.ts

CAP-026 · Command inbox — signal!'inference!'approval!'outcome flow

Full approval lifecycle with confidence score, evidence chain, policy displayed; all data seeded
status: Working Demo · module: artifacts/api-server/src/routes/approvals.ts, artifacts/lyte-command-center/src/pages/decision-center.tsx

CAP-028 · Action Debt, Decision Replay, Pressure Map, Ownership Drift (Lyte Command Center)

All four pages render with seeded operational data; no live data source connected
status: Working Demo · module: artifacts/lyte-command-center/src/pages/

OPEN RISKS & GAPS

No capabilities currently flagged as stub, broken, or undocumented for this product.

METHODOLOGY

Every figure in this brief is derived from artifacts/audit/platform-capability-manifest.json (manifest version 1.0.0). The readiness score weights live capabilities at 1.0, working demos at 0.75, and partials at 0.4. Capability statuses use SZL's published status definitions: live, working_demo, partial, stub, broken, undocumented. This brief is regenerated whenever the manifest changes — no manually edited claims.

